

Waza

Tw93
v3.28
MIT

Engineering habits you already know, turned into skills any AI agent can run.

8 engineering skills 6 agents supported 1 line to install 0 model lock-in

Models and tools come and go; the way good engineers actually work, planning, reviewing, debugging, reading the source, does not. Waza turns those habits into eight skills that run on whatever agent you have, so your workflow **never depends on a single model staying available**.

Plan & build


/think pressure-test a task into a plan
/design distinctive UI, never defaults
/learn map a new domain fast
/read distill any URL or PDF

Review & ship

/check review the diff with evidence
/hunt find a bug's true root cause
/write make prose sound human
/health audit your agent setup

Habits, not superpowers

The fashionable move is to bolt on a pack of "superpowers", dozens of skills promising an agent can suddenly do anything. But unconstrained capability is the problem, not the cure: it drifts into confident, generic, subtly wrong work. Waza adds no new powers. It encodes the habits a good engineer already has, so the model spends its capability on precision instead of breadth. Eight skills, not eighty, the few that actually decide quality, each carrying reference docs and real-failure gotchas rather than a promise that rots when a task gets specific.



Capability is the input, not the product. Constraints are what make it land.

Runs on Claude Code, Codex, Antigravity, OpenCode, Claude Desktop, and Pi. One workflow on every agent, installed once with `npm install tw93/Waza`.